

Solace at Baylor Scott & White — McKinney

Proposal for a no-cost, no-risk retrospective analysis of emergency department triage accuracy · Prepared July 2026 · Solace

Solace is a single AI layer for the emergency department encounter — intake, triage support, documentation and hand-off — designed to sit on top of Epic rather than replace anything. We are asking for one thing: permission to run a **retrospective analysis** of ED encounters that have already happened, and to give you the result.

What we are asking for

A 90-day retrospective study on de-identified historical ED data. **No patient contact. No clinical deployment. No change to any workflow. Nothing shown to any clinician.** We analyse encounters that are already closed and return a written finding.

What you get, whether or not you ever use our software


A measured answer to a question most emergency departments cannot currently answer about themselves:

Of the patients who arrived at McKinney and later required critical care, what acuity were they assigned at the door, and how many hours passed before the department recognised how sick they were?

That figure is your under-triage rate and your recognition delay. It is a patient-safety measure, it is not routinely computed by Epic reporting, and it belongs to you regardless of what happens next between us.

Why this matters — and why we cannot tell you your number yet

The largest study of triage accuracy in the United States reviewed 5,315,176 emergency department encounters across 21 hospitals. It found:



Separately, a 2026 analysis of 173,168 encounters found that 3.6% of patients deteriorated, and that **45% of those deteriorations occurred while the patient was still in the emergency department**. A related finding: only 40% of ordered vital signs were actually recorded.

We do not know McKinney's rate. The figures on the next page are published national rates applied to your published volume. They are an illustration of scale, not a claim about your department. Your actual rate may be materially better — McKinney is a Level III trauma centre and Baylor Scott & White publishes ED wait times well below national averages. *Establishing your real number is the entire purpose of the proposed analysis.*

Scale at McKinney

McKinney records **43,301 emergency department visits per year** — about 119 per day — across 192 beds, with Advanced Level III Trauma designation. Applying the published rates above to that volume gives the following *illustrative* figures.




Figure 1 — Scale of the question, at McKinney's volume.

The structural gap

Acuity is assigned once, at the door, by a triage nurse, in a few minutes, on limited information. It is then rarely revisited. A patient who is under-called at minute three carries that assignment for the remainder of a stay that may last many hours.





Figure 2 — The interval in which 45% of deteriorations occur.

Closing that gap prospectively is a clinical deployment and carries real risk, including alert fatigue. **We are not proposing that.** We are proposing to look backwards first, where the risk is zero, and to let the finding decide whether anything further is warranted.


What Solace is

One system covering the emergency department encounter end to end, rather than a single feature. Epic remains the system of record throughout; Solace reads from and writes back to it through standard interfaces (HL7 v2, FHIR R4, CDS Hooks).

Typical today



With Solace



69% of health system technology leaders name vendor management and integration as their top obstacle to executing AI. 62% want one comprehensive AI partner; 13% have one.

Figure 3 — Solace does not replace Epic. It replaces the need for several separate contracts around it.


Two design decisions we would want scrutinised

- **The model reports its own uncertainty.** Rather than a single acuity number, it returns a calibrated range with a stated coverage rate. A wide range is a signal to route to a clinician, not a prediction to act on.
- **Every output shows its basis.** Each estimate is accompanied by the specific inputs that drove it. This is deliberate: FDA's January 2026 Clinical Decision Support guidance turns on whether a clinician can independently review the basis for a recommendation, and we have built to that standard rather than around it.

Both decisions cost accuracy on paper. We think a model that declares when it is unsure is worth more in an emergency department than one that is confidently wrong 34% of the time, and we would rather defend that choice to you than hide it.

Proposed sequence

Three stages. Each one has to earn the next. We are asking only for the first.



We expect most conversations to stop after stage 1, and that is a legitimate outcome.

Figure 4 — Nothing reaches a patient until two prior stages have been reviewed by you.

What we would need from Baylor Scott & White

Item	Detail
A named clinical sponsor	An emergency physician or nursing leader willing to review the finding and tell us if it is useless. Roughly two hours per month.
Retrospective ED data	A de-identified or limited data set of closed encounters, under whatever agreement your privacy office requires — a Data Use Agreement, a BAA, or analysis conducted entirely inside your environment. We will work to whichever your compliance team prefers, including the option that no data leaves BSWH.
Definition review	Thirty minutes to agree how “under-triage” and “deterioration” should be defined for McKinney before we compute anything, so the result is one you recognise.
What we are not asking for	Money, an IT integration, EHR access, a workflow change, or any commitment beyond the 90 days.

What we would ask for if it works

Permission to proceed to stage 2, and — if the finding is genuinely useful — a reference we can show other Baylor Scott & White facilities. Nothing more at this stage.

Who we are, plainly

Solace is an early-stage team. We are not a funded company with hospital deployments behind us, and we would rather you learn that from us than from a reference check.

What exists today: a working system covering multilingual patient intake, a calibrated triage model with uncertainty reporting, ambient documentation, and interfaces to Epic, Oracle Health and athenahealth built on HL7 v2, FHIR R4 and CDS Hooks. It has a documented governance layer — model cards, consent enforcement, and an append-only audit log — because we built for a compliance review before we built for a demo.

What does not exist: a completed clinical validation, a health system deployment, or FDA clearance. Our triage model is currently trained on a public research dataset, not on real patient data. **That limitation is precisely why we are asking for a retrospective study rather than a clinical pilot.** A model validated against McKinney's own historical encounters is worth more — to you and to us — than any claim we could make today.

We are aware this is an unusual thing to put in a proposal. We think a Chief Medical Officer should be told what a system cannot yet do before being told what it might one day do.

Why McKinney, specifically

- **Volume sufficient to reach significance.** 43,301 annual visits gives enough events to produce a statistically meaningful finding within 90 days.
- **Level III trauma designation** means acuity assignment carries real consequence, and there is an existing quality infrastructure that already reviews it.
- **A high-performing department is the harder and more useful test.** Baylor Scott & White publishes ED wait times four to six times shorter than national averages. If a mistriage signal exists at McKinney, it exists everywhere. If it does not, that is a finding worth having and we would say so.

The regulatory context, briefly

- **CMS 2026 OPPS** adopted an Emergency Care Access and Timeliness measure, including the proportion of patients boarding beyond four hours. ED throughput moves from an operational concern to a reported one.
- **ONC HTI-1 §170.315(b)(11)** requires 31 source attributes to be surfaced for every predictive decision support intervention. Solace publishes these by default for its own models.
- **FDA CDS guidance (January 2026)** turns on whether a clinician can independently review the basis for a recommendation. Our outputs are built to that standard.

The single question we are asking

May we analyse McKinney's closed emergency department encounters, at no cost and with no patient contact, and tell you your own under-triage rate and recognition delay?

If the answer to that is yes, everything else can follow from what we find. If the finding is unremarkable, you will have a number you did not have before, and we will have learned something worth knowing.

References. Mistriage rates: Kaiser Permanente Division of Research, *JAMA Network Open*, n=5,315,176 — pubmed.ncbi.nlm.nih.gov/36930151. ED deterioration: *Annals of Emergency Medicine* (2026), n=173,168 — annemergmed.com. Vendor consolidation figures: health system technology leader survey reporting, 2026. CMS 2026 OPPS ED boarding measure — [ACEP](https://www.aacep.org). ONC HTI-1 Decision Support Interventions — [healthit.gov](https://www.healthit.gov). FDA Clinical Decision Support Software guidance, January 2026 — [fda.gov](https://www.fda.gov). McKinney facility data — [bswhealth.com](https://www.bswhealth.com).